

medical_costs_post

July 1, 2023

1 Get the Data

```
[52]: import numpy as np
import pandas as pd
import seaborn as sns
from scipy.stats import randint
import matplotlib.pyplot as plt
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.pipeline import make_pipeline
from sklearn.compose import ColumnTransformer
from sklearn.tree import DecisionTreeRegressor
from sklearn.metrics import mean_squared_error
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import OneHotEncoder
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import train_test_split
from sklearn.model_selection import RandomizedSearchCV

medical = pd.read_csv("insurance.csv")
```

2 Take a Quick Look at the Data Structure

```
[53]: medical.head()
```

```
[53]:
```

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520

```
[54]: medical.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1338 non-null   int64
1   sex         1338 non-null   object
2   bmi         1338 non-null   float64
3   children    1338 non-null   int64
4   smoker      1338 non-null   object
5   region      1338 non-null   object
6   charges     1338 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB

```

```
[55]: medical.region.value_counts()
```

```

[55]: southeast    364
      southwest    325
      northwest    325
      northeast    324
      Name: region, dtype: int64

```

```
[56]: medical.describe()
```

```

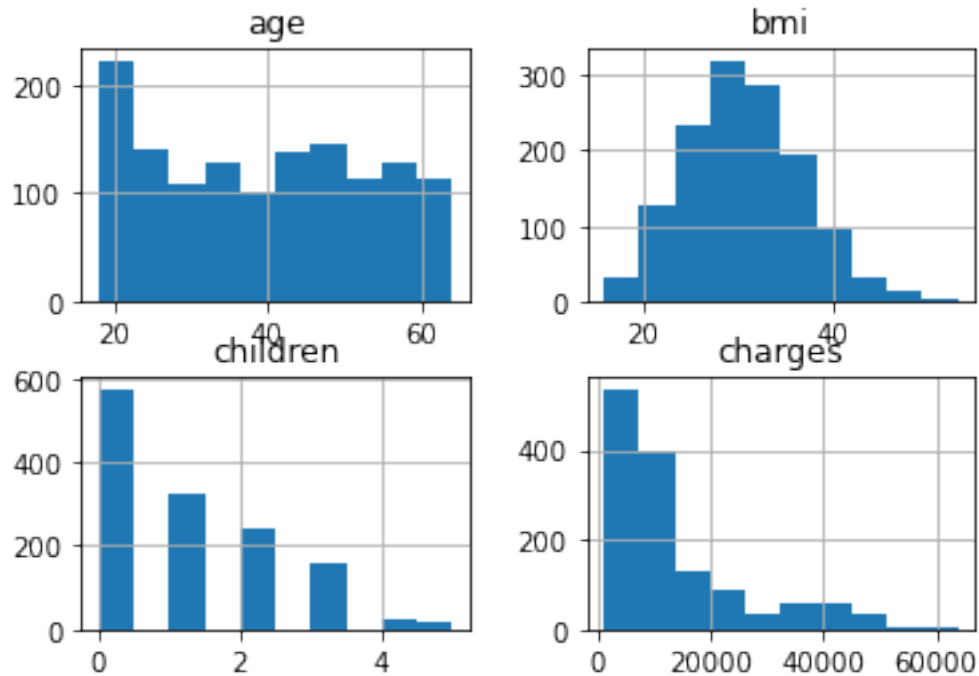
[56]:
      count      age      bmi      children      charges
count  1338.000000  1338.000000  1338.000000  1338.000000  1338.000000
mean     39.207025   30.663397    1.094918  13270.422265
std     14.049960    6.098187    1.205493  12110.011237
min     18.000000   15.960000    0.000000   1121.873900
25%     27.000000   26.296250    0.000000   4740.287150
50%     39.000000   30.400000    1.000000   9382.033000
75%     51.000000   34.693750    2.000000  16639.912515
max     64.000000   53.130000    5.000000  63770.428010

```

```

[57]: medical.hist()
      plt.show()

```



```
[58]: # age, children and charges are positively skewed
      # bmi is normal
```

2.1 Create a Test Set

```
[59]: RS = 13
```

```
[60]: X = medical.drop(["charges"], axis = 1)
      y = medical[["charges"]]
```

```
[61]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3,
      ↪ random_state = RS)
```

3 Explore and Visualize the Data to Gain Insights

```
[62]: X_train
```

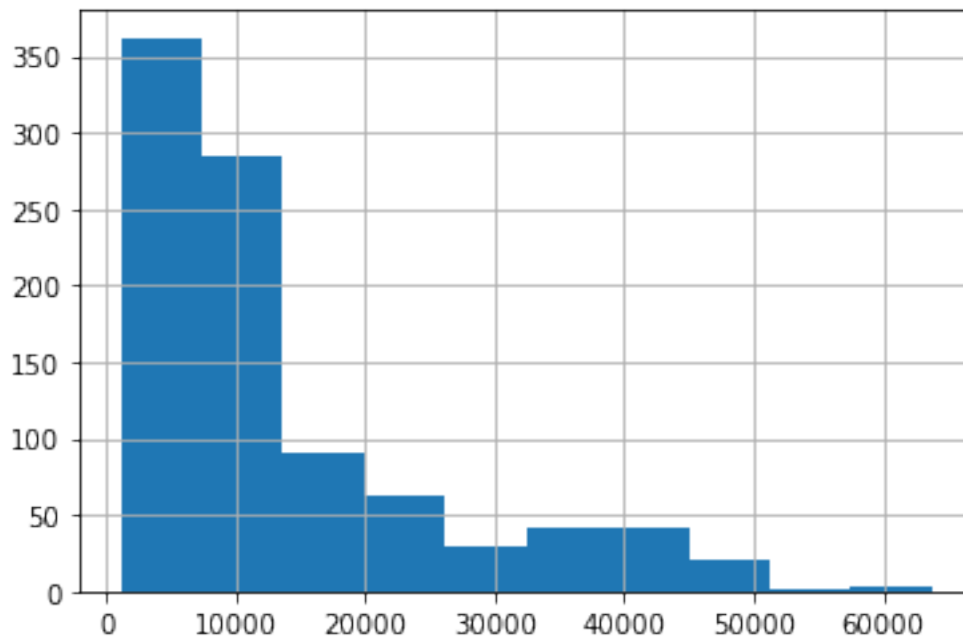
```
[62]:   age    sex    bmi  children  smoker    region
448   40  female  29.600         0     no  southwest
443   59  female  36.520         1     no  southeast
581   19   male  30.590         0     no  northwest
913   44  female  27.500         1     no  southwest
708   31  female  30.495         3     no  northeast
```

..
866	18	male	37.290	0	no	southeast
742	53	male	34.105	0	yes	northeast
74	44	male	27.400	2	no	southwest
176	38	male	27.835	2	no	northwest
338	50	male	32.300	1	yes	northeast

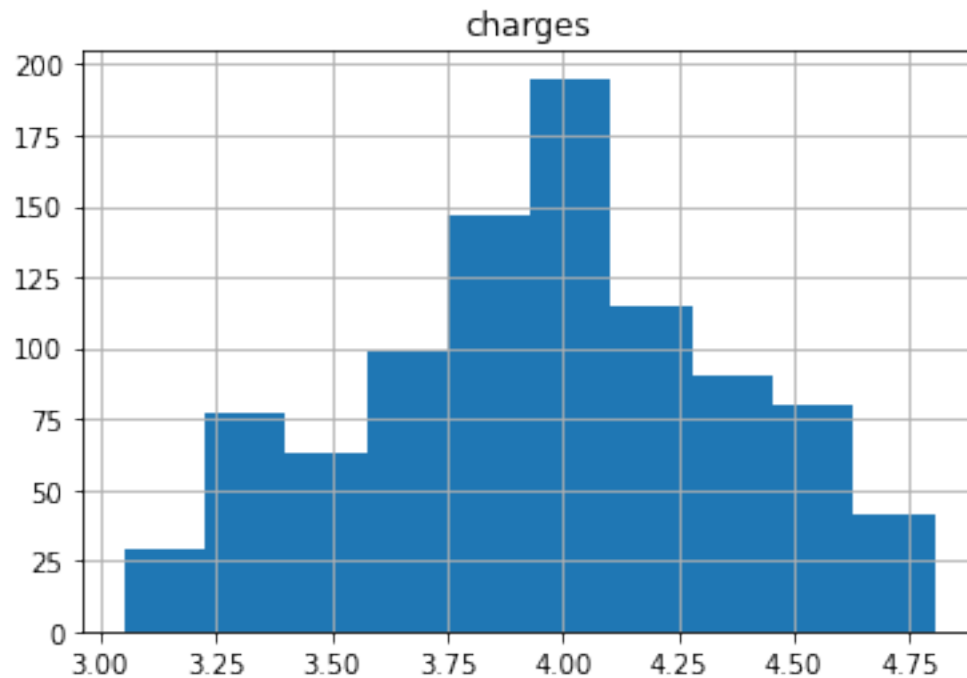
[936 rows x 6 columns]

```
[63]: train = X_train.copy()
      train["charges"] = y_train.copy()
```

```
[64]: train["charges"].hist()
      plt.show()
```



```
[65]: pd.DataFrame(np.log10(train.charges)).hist()
      plt.show()
```



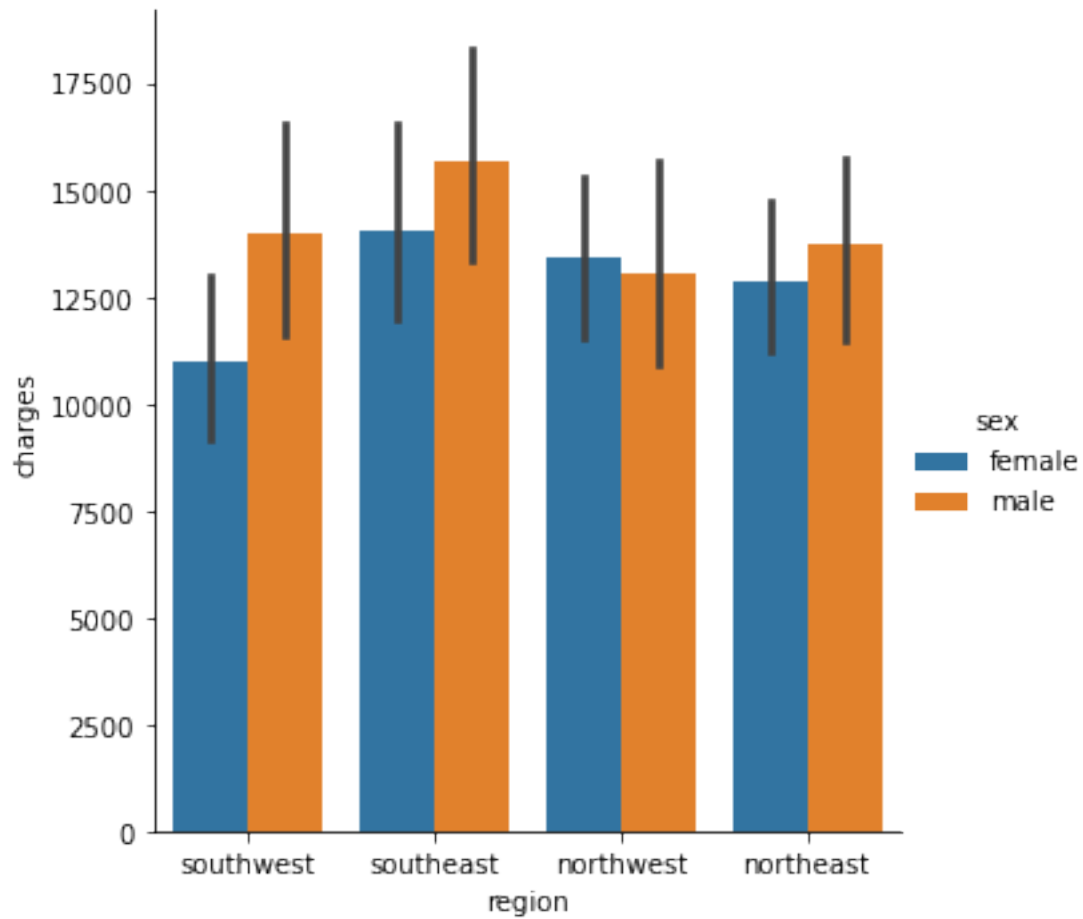
Now let's look at the mean charges by region

```
[66]: train[["region", "charges"]].groupby(["region"]).mean()
```

```
[66]:
```

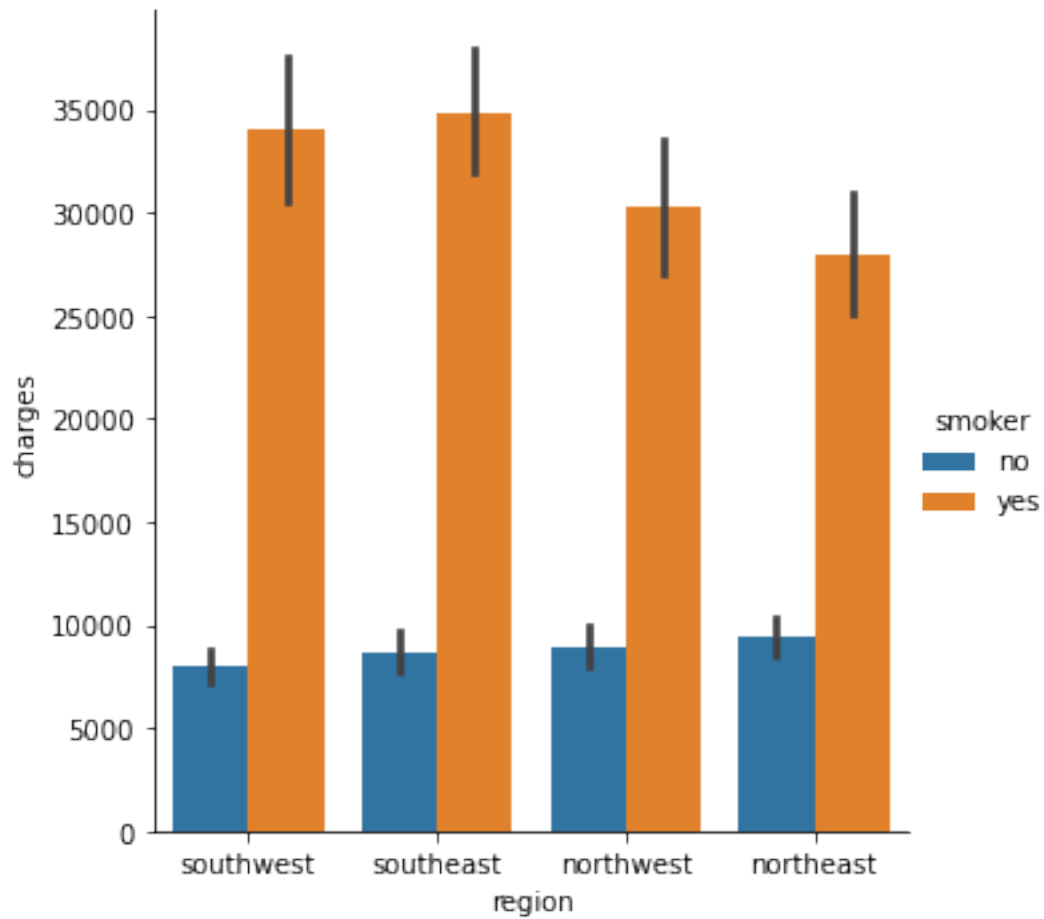
	charges
region	
northeast	13265.552265
northwest	13225.413686
southeast	14871.818962
southwest	12466.903493

```
[67]: sns.catplot(data = train, x = "region", y = "charges", hue = "sex", kind = "bar")
plt.show()
```



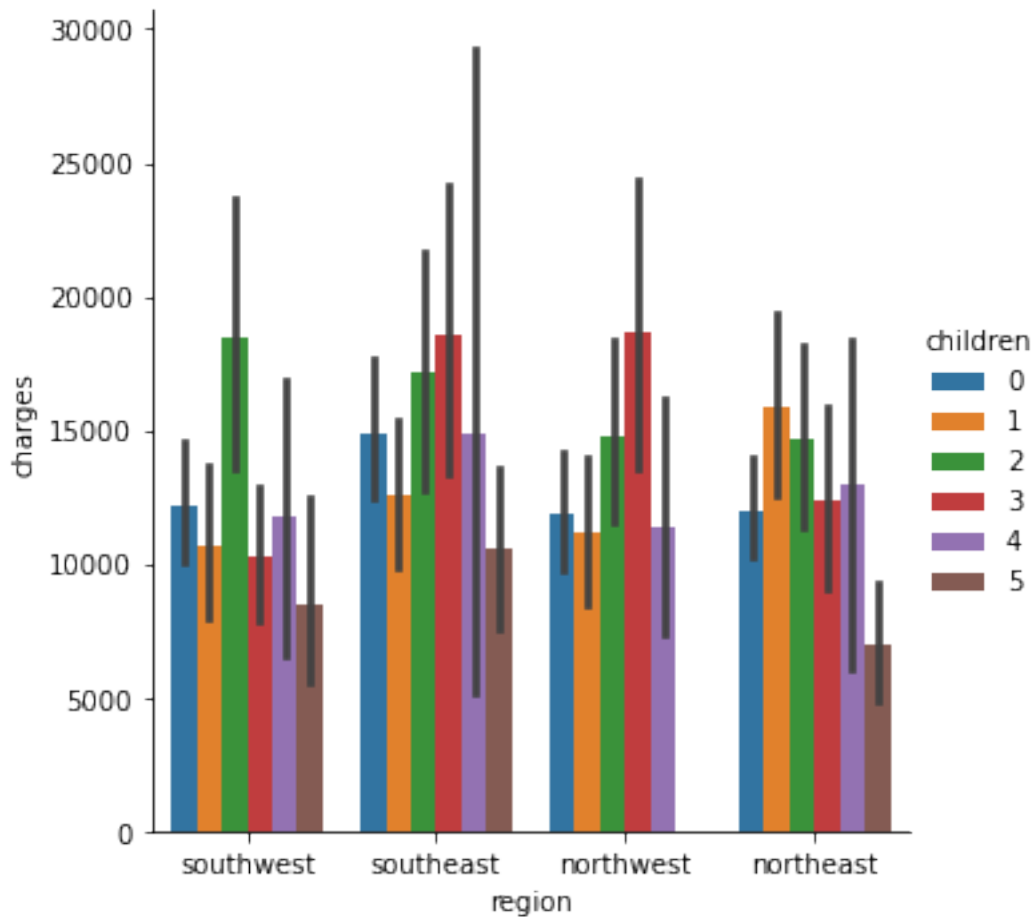
```
[68]: # males charges are more than females (in 3/4 cases)
      # southwest has the lowest charges for females
      # northwest has the lowest charges for males
      # southeast has the highest charges for both
```

```
[69]: sns.catplot(data = train, x = "region", y = "charges", hue = "smoker", kind = "
      ↪bar")
      plt.show()
```



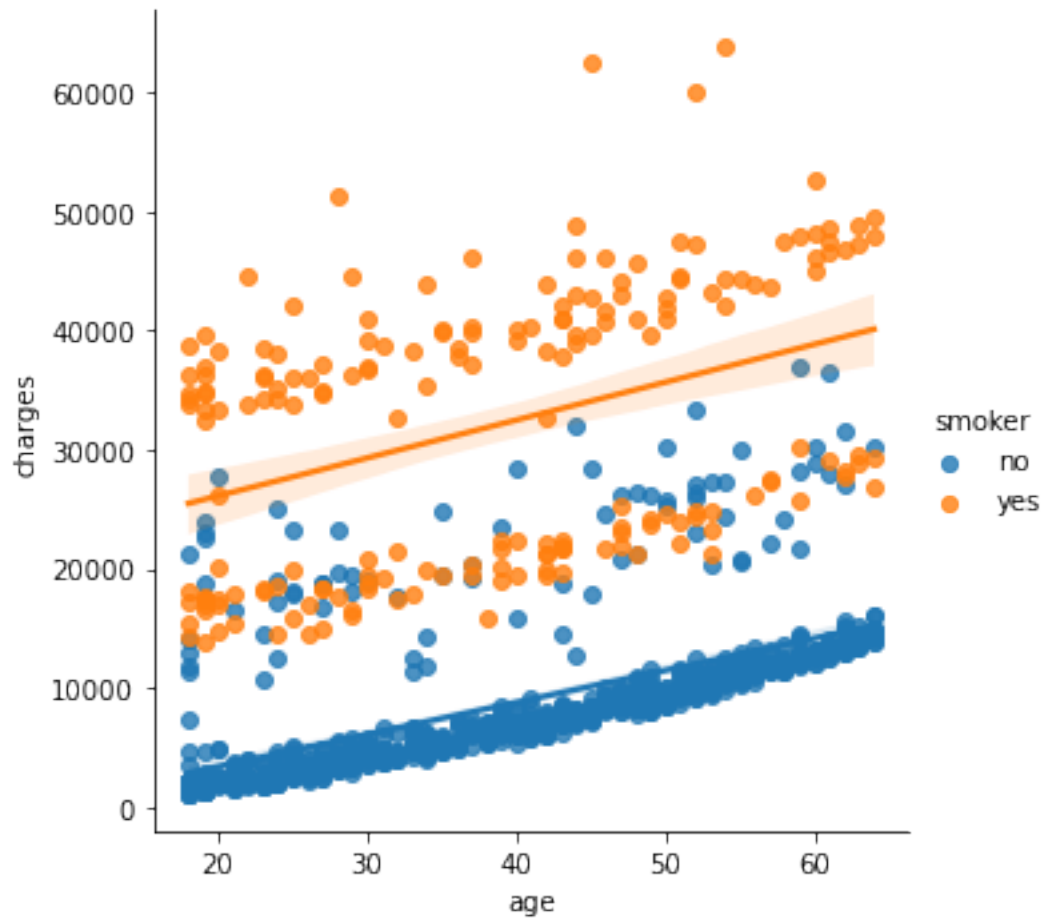
```
[70]: # smokers has 3 times more charges than non-smokers
      # north east has the highest charges for non-smokers, but the lowest charges
      ↪for smokers
```

```
[71]: sns.catplot(data = train, x = "region", y = "charges", hue = "children", kind =
      ↪"bar")
      plt.show()
```




```
[72]: # in southwest family with 2 children has the highest charge and with 5
      ↪ children have the lowest charge
      # in southeast family with 3 children has the highest charge and with 5
      ↪ children have the lowest charge
      # in northwest family with 3 children have the highest charge and with 1 child
      ↪ has the lowest charge
      # in northeast family with 1 child has the highest charge and with 5 children
      ↪ has the lowest charge
```

```
[73]: sns.lmplot(data = train, x = "age", y = "charges", hue = "smoker",)
      plt.show()
```

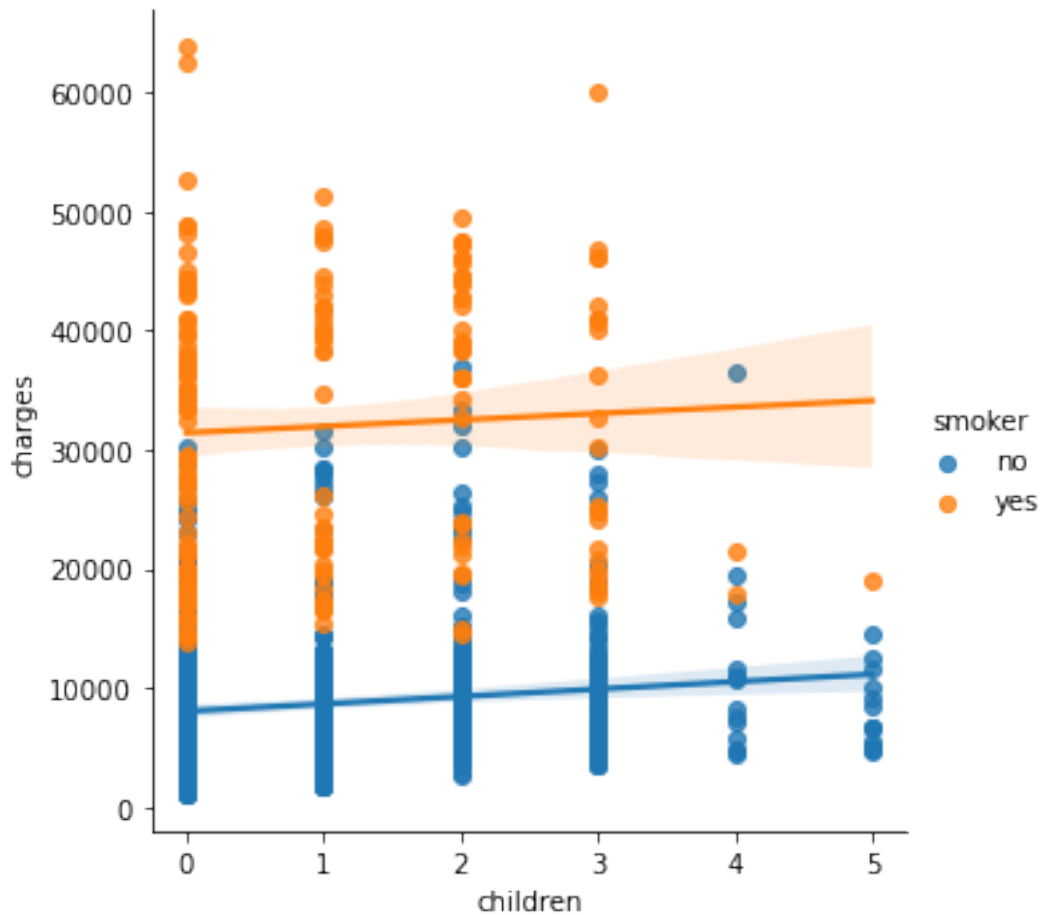
```
[74]: # age and smoking both positively effect on charges
```

```
[75]: sns.lmplot(data = train,x = "bmi", y = "charges", hue = "smoker",)  
plt.show()
```



```
[76]: # smoking highly effect on charges  
      # bmi effect on charges, but needs more clear plot for analyze
```

```
[77]: sns.lmplot(data = train,x = "children", y = "charges", hue = "smoker",)  
      plt.show()
```



```
[78]: # outliers of charges are mostly smokers and have less than 4 children
```

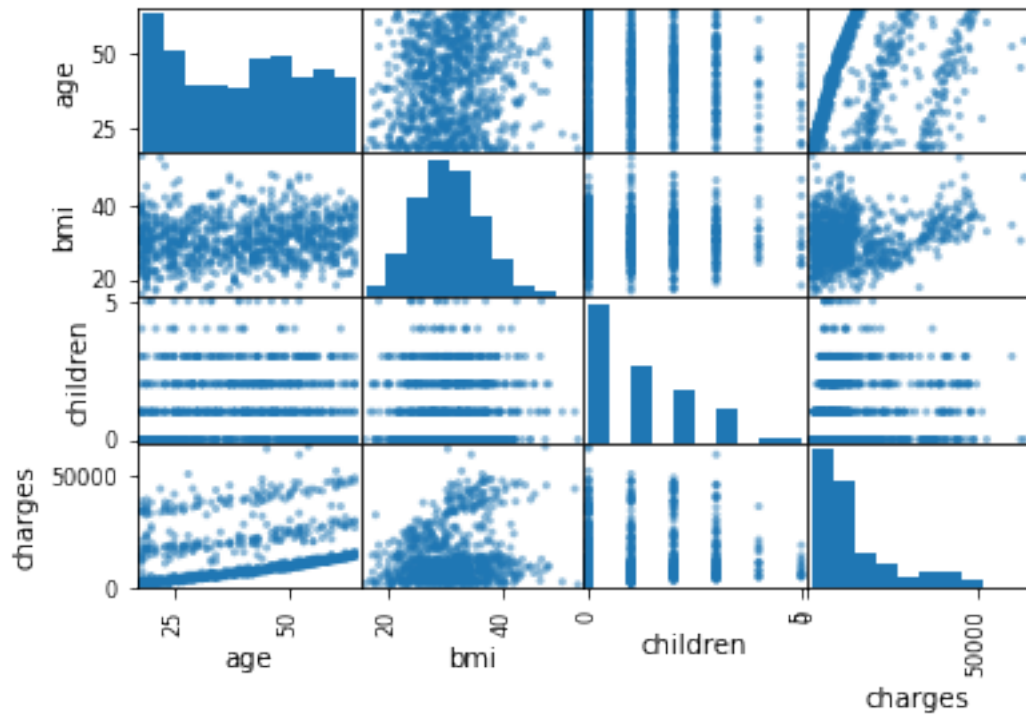
3.0.1 Look for Correlations

```
[79]: train.corr()
```


```
[79]:
```

	age	bmi	children	charges
age	1.000000	0.130788	0.025381	0.295932
bmi	0.130788	1.000000	0.023346	0.198454
children	0.025381	0.023346	1.000000	0.051307
charges	0.295932	0.198454	0.051307	1.000000

```
[80]: pd.plotting.scatter_matrix(train)
plt.show()
```



```
[81]: train.plot(kind = 'scatter', x = "bmi", y = 'charges', alpha = 0.3)
plt.show()
```



4 Transformation Pipelines

```
[82]: num_attribs = ["age", "bmi", "children"]
      cat_attribs = ["sex", "smoker", "region"]

      num_pipeline = make_pipeline(
          SimpleImputer(strategy="median"),
          StandardScaler())

      cat_pipeline = make_pipeline(
          SimpleImputer(strategy="most_frequent"),
          OneHotEncoder(handle_unknown="ignore", drop = 'first'))

      preprocessing = ColumnTransformer([
          ("num", num_pipeline, num_attribs),
          ("cat", cat_pipeline, cat_attribs)])

      X_train_perpared = preprocessing.fit_transform(X_train)

      print(X_train_perpared.shape)
      print(preprocessing.get_feature_names_out())
```

(936, 8)

['num__age' 'num__bmi' 'num__children' 'cat__sex_male' 'cat__smoker_yes'
 'cat__region_northwest' 'cat__region_southeast' 'cat__region_southwest']

5 Select and Train a Model

5.1 Train and Evaluate on the Training Set

```
[83]: lin_reg = Pipeline([
      ("preprocessing", preprocessing),
      ("linear_regression", LinearRegression())
  ])

      lin_reg.fit(X_train, y_train)
```

```
[83]: Pipeline(steps=[('preprocessing',
                        ColumnTransformer(transformers=[('num',
                                                           Pipeline(steps=[('simpleimputer',
                                                                 SimpleImputer(strategy='median')),
                                                                 ('standardscaler',
                                                                 StandardScaler())])),
                                                           ['age', 'bmi', 'children'])],
```

```

('cat',
Pipeline(steps=[('simpleimputer',
SimpleImputer(strategy='most_frequent')),
('onehotencoder',
OneHotEncoder(drop='first',
handle_unknown='ignore'))]),
['sex', 'smoker',
'region']]])),
('linear_regression', LinearRegression()))

```

```

[84]: y_predictions_lin = lin_reg.predict(X_test)
y_predictions_lin[:5]

```

```

[84]: array([[ 3578.74351013],
[10216.27622265],
[ 7398.59872957],
[ 6238.60129349],
[ 2562.75811151]])

```

```

[85]: lin_rmse = mean_squared_error(y_test, y_predictions_lin, squared=False)
lin_rmse

```

```

[85]: 5834.133007290974

```

```

[86]: dt_reg = Pipeline([
('preprocessing', preprocessing),
('decision_tree', DecisionTreeRegressor(random_state=RS))
])

dt_reg.fit(X_train, y_train)

```

```

[86]: Pipeline(steps=[('preprocessing',
ColumnTransformer(transformers=[('num',
Pipeline(steps=[('simpleimputer',
SimpleImputer(strategy='median')),
('standardscaler',
StandardScaler()))]),
['age', 'bmi', 'children']),
('cat',
Pipeline(steps=[('simpleimputer',
SimpleImputer(strategy='most_frequent')),
('onehotencoder',
OneHotEncoder(drop='first',
handle_unknown='ignore'))]),
['sex', 'smoker',
'region']]])),
('decision_tree', DecisionTreeRegressor(random_state=13))])

```

```
[87]: y_predictions_dt = dt_reg.predict(X_test)
      y_predictions_dt[:5]
```

```
[87]: array([ 2055.3249 , 12224.35085,  4320.41085,  5227.98875,  2457.21115])
```

```
[88]: dt_rmse = mean_squared_error(y_test, y_predictions_dt, squared=False)
      dt_rmse
```

```
[88]: 6122.432350426551
```

6 Better Evaluation Using Cross-Validation

```
[89]: tree_pipeline = Pipeline([
      ("preprocessing", preprocessing),
      ("tree", DecisionTreeRegressor(random_state = RS))
    ])

      tree_rmse = cross_val_score(tree_pipeline,
                                X_train, y_train,
                                scoring="neg_root_mean_squared_error",
                                cv=10)

      tree_rmse
```

```
[89]: array([-6265.29485138, -7377.28208902, -7085.38856846, -7997.65675418,
           -7567.24221623, -6805.94150941, -7849.43824253, -7307.06282526,
           -6137.19432622, -6172.97182115])
```

```
[90]: np.mean(tree_rmse * -1)
```

```
[90]: 7056.547320386965
```

```
[91]: np.median(tree_rmse * -1)
```

```
[91]: 7196.225696860822
```

```
[92]: rfr_reg = Pipeline([
      ("preprocessing", preprocessing),
      ("random_forest", RandomForestRegressor(random_state=RS))
    ])
      rfr_reg.fit(X_train, y_train)
```

```
/home/amyrmahdy/.local/lib/python3.10/site-packages/sklearn/pipeline.py:382:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples,), for example using
ravel().
```

```
self._final_estimator.fit(Xt, y, **fit_params_last_step)
```

```
[92]: Pipeline(steps=[('preprocessing',
                        ColumnTransformer(transformers=[('num',
                                                         Pipeline(steps=[('simpleimputer',
                                                         SimpleImputer(strategy='median')),
                                                         ('standardscaler',
                                                         StandardScaler())])),
                                                         ['age', 'bmi', 'children']),
                                                         ('cat',
                                                         Pipeline(steps=[('simpleimputer',
                                                         SimpleImputer(strategy='most_frequent')),
                                                         ('onehotencoder',
                                                         OneHotEncoder(drop='first',
                                                         handle_unknown='ignore'))])),
                                                         ['sex', 'smoker',
                                                         'region'])])),
                        ('random_forest', RandomForestRegressor(random_state=13))])
```

```
[93]: y_predictions_rfr = rfr_reg.predict(X_test)
      y_predictions_rfr[:5]
```

```
[93]: array([ 2401.83384026, 12397.9298855 ,  6870.0941465 ,  5123.836125 ,
              2875.93995009])
```

```
[94]: rfr_rmse = mean_squared_error(y_test, y_predictions_rfr, squared=False)
      rfr_rmse
```

```
[94]: 4251.558189926003
```

7 Fine-Tune Your Model

7.1 Randomized Search for Good Hyperparameters

```
[95]: full_pipeline = Pipeline([("preprocessing", preprocessing),
                                ("random_forest",
                                RandomForestRegressor(random_state=RS))
                                ])

param_distributions = {'random_forest__max_features': randint(low=2,high=20)}

rnd_search = RandomizedSearchCV(full_pipeline,
                                param_distributions=param_distributions,
                                n_iter=10,
                                cv=3,
                                scoring='neg_root_mean_squared_error',
                                random_state=RS)
rnd_search.fit(X_train, y_train.values.ravel())
```



```
[95]: RandomizedSearchCV(cv=3,
                        estimator=Pipeline(steps=[('preprocessing',
ColumnTransformer(transformers=[('num',
Pipeline(steps=[('simpleimputer',
                SimpleImputer(strategy='median')),
                ('standardscaler',
                StandardScaler())])),
['age',
'bmi',
'children']]),
('cat',
Pipeline(steps=[('simpleimputer',
                SimpleImputer(strategy='most_frequent')),
                ('onehotencoder',
                OneHotEncoder(drop='first',
                             handle_unknown='ignore'))])),
['sex',
'smoker',
'region']]))),
                        ('random_forest',
RandomForestRegressor(random_state=13))),
                        param_distributions={'random_forest__max_features':
<scipy.stats._distn_infrastructure.rv_frozen object at 0x7fd329ed8850>},
                        random_state=13, scoring='neg_root_mean_squared_error')
```

```
[96]: rn_res = pd.DataFrame(rnd_search.cv_results_)
rn_res.sort_values(by="mean_test_score", ascending=False, inplace=True)
rn_res.head(5)["mean_test_score"]
```

```
[96]: 4    -5039.948650
7    -5039.948650
6    -5093.283311
9    -5152.495411
0    -5213.865327
Name: mean_test_score, dtype: float64
```

```
[97]: final_model = rnd_search.best_estimator_ # includes preprocessing
feature_importances = final_model["random_forest"].feature_importances_
feature_importances
```

```
[97]: array([0.14900893, 0.17895356, 0.02726622, 0.00850256, 0.61334952,
0.00765398, 0.0084463 , 0.00681894])
```

```
[98]: sorted(zip(feature_importances, final_model["preprocessing"].
    ↳get_feature_names_out()),reverse=True)
```

```
[98]: [(0.6133495188155264, 'cat__smoker_yes'),
      (0.17895356150843594, 'num__bmi'),
      (0.14900892609866512, 'num__age'),
      (0.02726621957381335, 'num__children'),
      (0.008502558710902405, 'cat__sex_male'),
      (0.008446295284785842, 'cat__region_southeast'),
      (0.00765397739374062, 'cat__region_northwest'),
      (0.006818942614130391, 'cat__region_southwest')]
```

```
[99]: final_rfr_reg = Pipeline([
      ("preprocessing", preprocessing),
      ("final_random_forest", RandomForestRegressor(random_state=RS, max_features=
      ↪= 4))
  ])

  final_rfr_reg.fit(X_train, y_train.values.ravel())
```

```
[99]: Pipeline(steps=[('preprocessing',
                       ColumnTransformer(transformers=[('num',
                                                         Pipeline(steps=[('simpleimputer',
                                                         SimpleImputer(strategy='median')),
                                                         ('standardscaler',
                                                         StandardScaler()))],
                                                         ['age', 'bmi', 'children']),
                                                         ('cat',
                                                         Pipeline(steps=[('simpleimputer',
                                                         SimpleImputer(strategy='most_frequent')),
                                                         ('onehotencoder',
                                                         OneHotEncoder(drop='first',
                                                         handle_unknown='ignore'))],
                                                         ['sex', 'smoker',
                                                         'region'])])),
                       ('final_random_forest',
                       RandomForestRegressor(max_features=4, random_state=13))])
```

```
[100]: y_predictions_final_rfr = final_rfr_reg.predict(X_test)
       y_predictions_final_rfr[:5]
```

```
[100]: array([ 3251.5963769, 12139.4326819,  5544.7433706,  5053.354543 ,
              2925.5519552])
```

```
[101]: final_rfr_rmse = mean_squared_error(y_test, y_predictions_final_rfr,
      ↪squared=False)
       final_rfr_rmse
```

```
[101]: 4158.257085533594
```